


Measuring what a coding agent actually costs

Emil Lindfors | August 13, 2026



We have been experimenting on using an open source observability tool called OpenObserve for customers that want to keep telemetry on prem, and thus I wanted to see how to integrate LLM telemetry from an agent harness there.

The agent harness is the Codex fork from [the previous post](#), and the cost claims in [the first post](#) come out of this work. Getting the numbers right took considerably longer than the fork did.

Everything here is self-hosted and open source. No vendor SaaS, no per-seat pricing, and no prompt content shipped to a third party. That last one you have to actively prevent, and there is a section on it below.

The stack

Component	Role	Version
OpenObserve	Store, query, dashboards for logs, metrics and traces	latest, podman container


Component	Role	Version
OpenTelemetry Collector (otelcol-contrib)	Ingest, transform, route	0.158.0
Keycloak	OIDC SSO in front of OpenObserve	realm lindfors
nginx	TLS and an authenticated public OTLP endpoint	—
Codex CLI	LLM client, emits its own OTLP metrics and traces	0.1.1
OpenRouter Broadcast	LLM gateway, ships spans with token counts and cost	—

OpenObserve Enterprise is free under 200 GB/day of ingest. One host does not get anywhere near that.

They do have built in telemetry for a lot of these tools now so it's even easier than what I made it out to be.

So if you start today, check OpenObserve's own integrations first. What follows is the collector route, and you still want the collector for the second half of this post, which is about dropping noise before it costs storage.

Three ingest paths, one collector, one store:



OpenObserve speaks OTLP directly, so strictly the collector is optional. It is there for two jobs. It routes to per-destination streams (OpenObserve picks the stream from a `stream-name` header, so you need one exporter per stream, otherwise everything lands in `default` and retention becomes unmanageable), and it drops noise before it costs storage. That second job turned out to matter more than anything else here.

The OTLP receiver binds to loopback and has no auth of its own. Anything from off-host arrives through nginx, which terminates TLS, enforces basic auth, and strips the credential before proxying so the collector never sees it:

```
location /otlp/ {
    auth_basic "OTLP ingest";
    auth_basic_user_file /etc/nginx/otlp.htpasswd;

    proxy_pass http://127.0.0.1:4318/;

    # do not leak the ingest credential to the collector
    proxy_set_header Authorization "";

    client_max_body_size 32m;
}
```

`client_max_body_size 32m` is not decoration. OTLP batches carrying full prompt text get large and the 1 MB default silently 413s them.

Getting the numbers, three times

I got the numbers wrong twice before they were right. Here are all three attempts, because the two wrong ones looked fine on a chart.

Attempt 1: derive metrics from OpenRouter spans

OpenRouter Broadcast emits a span per generation with token counts and cost as attributes. You can query the trace stream directly, but that binds cost history to trace retention and gets slow over long ranges. So: convert spans into real metric time series with the collector's `sum` connector.

```
connectors:
  sum:
    spans:
      gen_ai.usage.input_tokens:
        source_attribute: gen_ai.usage.input_tokens
        conditions:
          - 'attributes["gen_ai.usage.input_tokens"] != nil'
        attributes:
          - key: gen_ai.request.model
```

Two traps in that block, both of which I only found by attaching the debug exporter and reading raw datapoints.

The emitted value is multiplied by the number of `attributes` entries. Listing two keys turned a 500-token span into 1000. Keep exactly one.

A span missing *any* listed attribute is dropped silently. No datapoint, no warning, no metric. Listing `gen_ai.system` alongside the model would have discarded every span that lacked it, and I'd have had a chart that looked fine.

A third one came from the `hostmetrics` side. The `sum` connector emits Delta temporality, so `sum()` over a time bucket is correct as written. Other sources emit **cumulative** counters, where each datapoint is a running total, and summing those per hour double-counts every snapshot. An hour of about 1.08 M input tokens came out as roughly 1.6 M.

```
processors:
  cumulative_to_delta:
    include:
      metrics: [gen_ai.usage.input_tokens, gen_ai.usage.output_tokens,
gen_ai.usage.total_cost]
      match_type: strict
```

Find out which temporality your source emits before writing a single aggregation. None of these three produced an error. They produced numbers that were plausibly wrong, and nothing flagged them.

Attempt 2: use Codex's own metrics

Codex CLI exports its own OTLP metrics, and they look like the right source: first-party, no span-derivation fragility, about 35 metric families covering turns, tool calls, latency, SSE events, sqlite init, skills selection.

The shape differs from the OTel `gen_ai` semantic conventions, so a few notes if you go this way. All token counts live in one histogram, `codex_turn_token_usage`, split by a `token_type` label (`input`, `cached_input`, `cache_write_input`, `output`, `reasoning_output`, `total`), not one metric per direction; the `_sum` stream carries the count and `_count` is the number of turns. `token_type = 'total'` already includes input plus output, so any per-model panel has to filter to a single type or it double-counts. And the model label is plain `model`, not `gen_ai.request.model`.

Two things killed it as the primary source.

Codex reports no cost metric at all. There is no `codex_*` equivalent of `gen_ai.usage.total_cost`. Token counts are there, so you could reconstruct cost from your own price table, but nothing on the wire carries it.

And the token metrics report zero. `codex_turn_token_usage` records `0` for every `token_type`, including `input` and `total`, across heavy multi-hour sessions — all 16 histogram buckets land in `le=0`. The same zeros show up in Codex's own `session_task.turn` spans, which rules out the collector and the store. It's Codex's accounting.

Meanwhile the latency and activity metrics from the same source are fine and report real numbers. A telemetry source can be half-broken, and on a dashboard the broken half looks exactly like “no usage yet”.

Attempt 3: each source for what it's good at

What stuck:

Signal	Source	Why
Cost (total, input, output, unit prices)	OpenRouter <code>LLM Generation</code> spans	Only source that has it
Tokens (input, output, cached)	OpenRouter <code>LLM Generation</code> spans	Codex's are zero
Requests, turns	Codex metrics	Client-side truth, including failures
Latency (turn end-to-end, TTFT)	Codex metrics	Measured where the user waits

The trap on the OpenRouter side is that a single generation produces a span tree, and only the root carries real numbers:

LLM Generation	inp=92495	out=535	total_cost=0.001859	←
the real one				
└ generation	inp=0	out=0	cost=-0.0	
└ provider attempt 1: DeepInfra	inp=0	out=0	cost=-0.0	

Filter on `operation_name = 'LLM Generation'` or every number is zero. I nearly built the whole dashboard on a child span.

Cost fields also arrive as strings. `gen_ai_usage_total_cost` is `"0.001859058"`, not a float, so every panel needs `CAST(... AS DOUBLE)`. Token counts are numeric. There's no consistency to it, it's just a thing to know.

Trying to join the two sources

The obvious next question is whether you can attach a cost figure to a specific Codex turn, or to a project. As far as I can tell, not without changing something upstream.

What I checked:

- **Trace context.** Zero shared `trace_id`s. Either Codex doesn't inject `traceparent` into the OpenRouter call, or OpenRouter doesn't honour it.
- **A client-supplied id.** OpenRouter surfaces `trace_metadata.openrouter.user_id` from the `user` field of the request body. Codex sends nothing there.

- **Wall clock.** The only real link. Codex's `chat.stream_request` trails OpenRouter's `LLM Generation` by a strikingly stable 227–242 ms, with 5–48 s between calls, so pairing them by eye is unambiguous.

That last one looks like a solution. I abandoned it for two reasons.

The first is that OpenObserve's DataFusion silently returns zero rows for inequality joins. An interval join (`ON o.start_time BETWEEN c.start_time AND c.end_time`) returns an empty result set. Not an error, not a warning, just nothing. Equi-joins are fine. `UNION ALL` and cross joins panic outright:

```
assertion failed: variadic_counts.is_empty()
```

So the only workable shape is an equi-join on a bucketed timestamp, which means hard-coding a 235 ms offset and hoping it never drifts across a collector restart, a network change or an NTP correction. That is not something I want behind a number I show a client.

What works instead, for free: a separate OpenRouter API key per project. The key name arrives as `trace_metadata.openrouter.api_key_name` on every span, which gives exact project-level attribution with no join at all. It's coarser than per-turn and it's correct, and I'll take correct over precise.

If you want to correlate two telemetry sources, design the join key in from the start.

We have thought about using a split thing like using both OpenRouter plus Codex telemetry to get the full picture but that is a bigger project that we haven't started yet.

The dashboard below is the small version of that: both sources on one page, read side by side.

Agent traces are enormous

Codex with tracing on emitted **118,469 spans in one working session**. 118 of them were about the LLM. The rest is Rust async runtime internals:

```
FramedRead::poll_next      1626
poll_ready                  1626
try_reclaim_frame           1527
FramedRead::decode_frame    864
FramedWrite::flush          838
pop_frame                   825
...
```

0.1% signal. h2 frame handling and tokio worker polls, faithfully shipped across the network and stored forever. This is the biggest operational problem in the whole setup, bigger than any of the metric bugs, because it's pure cost with no upside at all.

The store does the same thing to itself. OpenObserve at `INFO` logs every ingest and every search. The collector tails container logs and ships them to OpenObserve, which logs that, which gets shipped back. About 10k lines an hour of pure self-reference. Two settings are needed to stop it, because `ZO_LOG_LEVEL` on its own does nothing — the Rust `tracing` subscriber honours `RUST_LOG`, and that's the one that takes effect:

```
ZO_LOG_LEVEL=warn
RUST_LOG=warn
```

There's a third variant neither setting touches. OpenObserve dumps DataFusion physical query plans to stdout during compaction, printed as raw table output rather than through the tracing subscriber, so no config flag suppresses them. Those get dropped at the collector:

```
processors:
  filter/plan_dumps:
    error_mode: ignore
    logs:
      log_record:
        - 'IsMatch(body, "(DataSourceExec|ProjectionExec|AggregateExec|SortExec|FilterExec)")'
```

Budget for noise suppression as a real part of the work. Every component here over-reports by default and each one needed a different mechanism to quiet down: an env var, a second env var, a collector filter, and a client-side sampling policy I still haven't written.

Your prompts end up in your logs

By default OpenRouter Broadcast ships `gen_ai.input_messages` and `gen_ai.output_messages`. That is the entire conversation, system prompt included. On this setup it means 90k-token prompts sitting in the trace store. It's a privacy question and it's also most of the storage.

Turning on LLM tracing means logging your prompts unless you take a specific step to stop it. There's a provider-side toggle; the collector-side attribute drop is the belt-and-braces version and it covers every source at once:

```
processors:
  transform/strip_prompt_content:
    error_mode: ignore
    trace_statements:
      - context: span
```

```
statements:
  - delete_key(attributes, "gen_ai.input_messages")
  - delete_key(attributes, "gen_ai.output_messages")
  - delete_key(attributes, "gen_ai.system_instructions")
```

I have written it down and not yet applied it on this host. It is in the open items below.

The dashboard

One dashboard, `LLM Usage`, ten panels, both sources on a shared hourly axis and read together by eye, not joined.

Panel	Source
Cost (USD) per hour by model	OpenRouter traces
Total cost by model	OpenRouter traces
Cost by upstream provider	OpenRouter traces
Total tokens by model	OpenRouter traces
Tokens per hour by model	OpenRouter traces
Cached input tokens per hour	OpenRouter traces
API requests per hour by model	Codex metrics
Conversation turns by model	Codex metrics
Slowest turn end-to-end (ms)	Codex metrics
Slowest time to first token (ms)	Codex metrics

A representative panel, with both the `CAST` and the root-span filter:

```
SELECT histogram(_timestamp, '1 hour')           as "x_axis_1",
       sum(CAST(gen_ai_usage_total_cost AS DOUBLE)) as "y_axis_1",
       gen_ai_request_model                       as "breakdown_1"
FROM   "default"
WHERE  operation_name = 'LLM Generation'
GROUP BY "x_axis_1", "breakdown_1"
ORDER BY "x_axis_1"
```

Dashboards are managed over the REST API, not the UI, so they are reviewable and re-creatable:

```
GET  /api/{org}/dashboards
PUT  /api/{org}/dashboards/{id}?folder=default&hash={hash}
```

The `hash` is a concurrency token from the `GET`. Stale hash, rejected write.

Sample numbers over about two days of real use: **\$0.184 across 5.97 M tokens, of which 2.0 M were prompt-cache hits**. The cache panel is the most actionable thing on the

board, since it's the difference between an eighteen-cent day and a much worse one. That figure is the evidence behind the cost claims in the [business post](#): at this scale tooling cost is noise, and now there is a number behind the sentence.

Learning by using

Here is the lens for this one. Nathan Rosenberg's *Inside the Black Box* (1982) has a chapter called "Learning by using", and its argument is simple. For a complicated piece of equipment, you cannot know what it really costs to run, or what it is really good at, from its design or its price list. You find out by using it for a long time and paying attention. His examples are aircraft: the airlines learned the maintenance costs, the fuel economy and the right routes for a new airframe over years of service, and a great deal of that knowledge fed back into stretched versions of the same plane. He splits the learning in two. *Embodied* learning changes the design. *Disembodied* learning changes how you operate the thing you already have, and he found that the second kind often moved the economics more.

Rosenberg wrote about jet engines and machine tools, and an agent harness is software, so the transfer is not automatic. I think it holds, because a coding agent behaves much more like a machine tool than like a subscription: what it costs depends on how you run it. Read the two months above through the lens and every finding is a learning-by-using finding:

- **The price list told me nothing.** \$0.09 per million tokens is design information. The number that mattered, 2.0 M of 5.97 M tokens served from cache, only exists after two days of real use with a meter attached.
- **The disembodied learning is where the money is.** Prompt caching, reasoning effort, and which model gets which task are all operating practice. None of them required a new model or a patch, and together they move cost-per-task more than switching vendor would.
- **The cost of measuring is itself learned by using.** Nobody's documentation says that a tracing session produces 118,469 spans. You find out when the disk fills.
- **The instrument has to be there before the learning can start.** Rosenberg's airlines kept maintenance logs. If you are choosing an agent or a model on price, you are choosing on the information that use will overturn first, so put the meter in before you decide.

Access control, briefly

OpenObserve's Dex integration points straight at Keycloak's OIDC endpoints, with no Dex container involved. Four things that cost me time:

`O2_DEX_BASE_URL` must be the realm root, because OpenObserve fetches `BASE_URL + /.well-known/openid-configuration` and Keycloak serves that at the realm root, not under `/protocol/openid-connect`. The Keycloak client needs an audience mapper adding `openobserve` to the access token's `aud`, since Keycloak defaults it to `account`, which fails validation. Org and role come from per-user Keycloak attributes (`o2_org` / `o2_role`) surfaced as claims by attribute mappers, deliberately not the `groups` claim, whose value is `admin` and would put users in a non-existent `admin` org while all data ingests into `default`; those attributes also have to be declared in the realm's user profile or Keycloak drops them without saying so. And roles of external users can't be edited in OpenObserve at all, because they're read from these claims on every login.

Container logs get their service name from the filename, which means starting containers with an explicit log path:

```
podman run --log-driver k8s-file --log-opt path=/var/log/containers/<name>.log
```

The default podman path exposes only a container ID. An earlier version of my config hardcoded `source: keycloak` onto every container's logs, which was simply wrong — Keycloak's log is about 4 KB where OpenObserve's is 1.4 MB over the same window, so the mistake was easy to spot once I looked.

Still open

- **Codex trace sampling.** 99.9% of spans are runtime internals. Needs a client-side filter, or a `filter` processor on `service.name IN (codex_*)` that drops everything but the handful of meaningful operations.
- **Stripping prompt content** at the collector, per above.
- **Retention.** `/opt/openobserve/data` is 238 MB and growing, and retention isn't configured, so it grows without bound.
- **Codex token metrics reporting zero**, which I should report upstream.
- The `sum` connector and its `cumulative_to_delta` entry are dead weight now that cost is read from spans directly, and I haven't removed them.
- **Whether OpenRouter stays in the stack at all.** The cost and token numbers here all come through it, and the gateway question is open again:

We are still undecided if OpenRouter is the way to go now that they have been bought by Stripe and enterprise functionality like EU processing starts at 25k USD a month.

If it goes, the span source goes with it, and the split-source project above becomes the way to get cost back.

If you build this yourself

Attach the debug exporter before anything else. Every metric bug above was invisible in the dashboard and obvious in raw datapoints.

Then, in the order they bit me:

- Check temporality before writing any aggregation. Cumulative summed per bucket is the easiest way to be 50% wrong.
- Find the root span. Span trees put the real numbers in one place and zeros everywhere else.
- Design the join key in, or don't join.
- Verify a query returns rows before trusting an empty result. A silent zero-row join looks exactly like "no data yet".
- Assume every component over-reports, and that each one needs a different off switch.
- Decide about prompt content on day one, not after 90k-token prompts are already on disk.

The stack is running, the dashboard is up, and the next thing that changes it is the gateway decision. Until then, eighteen cents a day.